

```

import numpy as np
import pdb

class KNN(object):

    def __init__(self):
        pass

    def train(self, X, y):
        """
        Inputs:
        - X is a numpy array of size (num_examples, D)
        - y is a numpy array of size (num_examples, )
        """
        self.X_train = X
        self.y_train = y

    def compute_distances(self, X, norm=None):
        """
        Compute the distance between each test point in X and each
        training point
        in self.X_train.

        Inputs:
        - X: A numpy array of shape (num_test, D) containing test data.
        - norm: the function with which the norm is taken.

        Returns:
        - dists: A numpy array of shape (num_test, num_train) where
        dists[i, j]
        is the Euclidean distance between the ith test point and the jth
        training
        point.
        """
        if norm is None:
            norm = lambda x: np.sqrt(np.sum(x**2))
            #norm = 2

        num_test = X.shape[0]
        num_train = self.X_train.shape[0]
        dists = np.zeros((num_test, num_train))
        for i in np.arange(num_test):

            for j in np.arange(num_train):
                #
                ===== #
                # YOUR CODE HERE:
                #   Compute the distance between the ith test point and the
                jth

```

```

        # training point using norm(), and store the result in
dists[i, j].
        #
===== #
        # X is x-test
        dists[i, j] = norm(X[i] - self.X_train[j])

        #
===== #
        # END YOUR CODE HERE
        #
===== #

    return dists

def compute_L2_distances_vectorized(self, X):
    """
    Compute the distance between each test point in X and each
training point
    in self.X_train WITHOUT using any for loops.

    Inputs:
    - X: A numpy array of shape (num_test, D) containing test data.

    Returns:
    - dists: A numpy array of shape (num_test, num_train) where
dists[i, j]
        is the Euclidean distance between the ith test point and the jth
training
        point.
    """
    num_test = X.shape[0]
    num_train = self.X_train.shape[0]
    dists = np.zeros((num_test, num_train))

    # =====
#
# YOUR CODE HERE:
# Compute the L2 distance between the ith test point and the jth
# training point and store the result in dists[i, j]. You may
# NOT use a for loop (or list comprehension). You may only use
# numpy operations.
#
# HINT: use broadcasting. If you have a shape (N,1) array and
# a shape (M,) array, adding them together produces a shape (N,
M)
# array.
# =====
#

```

```

        # broadcasting: kind of like amplifying the array

        # axis = 1 means sum the cols
        # reshape(-1, 1): (-1) takes the shape of the last dimension, puts
1 as that dimension
        # (2, ) -> (2, 1)

        dists = np.sqrt(np.sum((self.X_train)**2, axis=1) + np.sum((X)**2,
axis=1, keepdims=True) - 2*X.dot(self.X_train.T))

    pass

    # =====
#
    # END YOUR CODE HERE
    # =====
#

    return dists

def predict_labels(self, dists, k=1):
    """
    Given a matrix of distances between test points and training
    points,
    predict a label for each test point.

    Inputs:
    - dists: A numpy array of shape (num_test, num_train) where
dists[i, j]
    gives the distance between the ith test point and the jth
    training point.

    Returns:
    - y: A numpy array of shape (num_test,) containing predicted
    labels for the
    test data, where y[i] is the predicted label for the test point
    X[i].
    """
    num_test = dists.shape[0]
    y_pred = np.zeros(num_test)
    for i in np.arange(num_test):
        # A list of length k storing the labels of the k nearest
neighbors to
        # the ith test point.
        closest_y = []
        #
    ===== #

```

```

        # YOUR CODE HERE:
        # Use the distances to calculate and then store the labels of
        # the k-nearest neighbors to the ith test point. The function
        # numpy.argsort may be useful.
        #
        # After doing this, find the most common label of the k-
nearest
        # neighbors. Store the predicted label of the ith training
example
        # as y_pred[i]. Break ties by choosing the smaller label.
        #
===== #
        sorted_i = np.argsort(dists[i]) # sort the distances, the array
elements are the j's
        closest_indices = sorted_i[:k] # take only the first k elements
- but how to change the j to the label? ???

        # print(closest_indices)
        closest_y = self.y_train[closest_indices]
        # for j in closest_indices:
        #     labels.append(self.y_train[j])
        # print("i got here")
        y_pred[i] = np.bincount(closest_y).argmax() # save the most
frequent element

        pass

        #
===== #
        # END YOUR CODE HERE
        #
===== #

    return y_pred

```