

```

import numpy as np

class Softmax(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001

    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on
        minibatches of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i]
        = c means
          that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """

        # Initialize the loss to zero.
        loss = 0.0

        # =====
        #
        # YOUR CODE HERE:
        #   Calculate the normalized softmax loss. Store it as the
        variable loss.
        #   (That is, calculate the sum of the losses of all the training
        #   set margins, and then normalize the loss by the number of
        #   training examples.)
        # =====
        #

        # print(X.shape)

```

```

    # print(self.W.T.shape)
    ex_sc = 0

    for i in range(X.shape[0]):
        e_x = np.exp(np.matmul(self.W, X[i])) # exp(wT * X) <-- exp of
score of each class

        ex_sc = e_x # sum of softmax scores, for the denominator

        ex_sc = ex_sc/(np.sum(ex_sc)) #

        log_sc = np.log(ex_sc) # take the log likelihood

        loss = loss - log_sc[y[i]]

    loss = loss/(X.shape[0])

    # =====
#
    # END YOUR CODE HERE
    # =====
#

    return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.

    Output: grad -- a matrix of the same dimensions as W containing
    the gradient of the loss with respect to W.
    """

    # Initialize the loss and gradient to zero.
    loss = 0.0
    grad = np.zeros_like(self.W)

    # =====
#
    # YOUR CODE HERE:
    # Calculate the softmax loss and the gradient. Store the
gradient
    # as the variable grad.
    # =====
#

    # ex_sc = 0
    # grad_sum = 0
    # smax_all = [X.shape[0]][self.W.shape[0]] # n x c
    # for i in range(X.shape[0]):
    #     e_x = np.exp(np.matmul(self.W, X[i])) # exp(wT * X) <-- exp of

```

```

score of each class
    #   ex_sc = e_x/(np.sum(e_x)) # normalize the scores - becomes row
of softmax
    #   smax_all[i] = ex_sc

    #   ex_sc += e_x # sum of softmax scores, for the denominator

    # print("x shape", X.shape) # 500
    # print("w shape", self.W.shape) # 10

    for i in range (X.shape[0]): # i is range 500
        e_x = np.exp(np.matmul(self.W, X[i])) # exp(wT * X) <-- exp of
score of each class (numerator)
        # print("e_x shape", e_x.shape[0]) # 500
        ex_sum = e_x/np.sum(e_x, axis=0) # ex/sum ex (so this adds in
the denominator)
        # print(np.shape(ex_sum))
        log_sc = np.log(ex_sum) # take the log likelihood

        loss = loss - log_sc[y[i]]

    for j in range (self.W.shape[0]): # j is range 20
        if (j == y[i]):
            grad[j] += (ex_sum[j]-1)*X[i]
        else:
            grad[j] += (ex_sum[j])*X[i]

    # print("ex sum", ex_sum)
    # print("loss before normalization", loss)
    # print("log sc", log_sc)
    grad = grad/(X.shape[0]) # normalize the grad
    loss = loss/(X.shape[0])
    print("for loop loss", loss)
    # print(grad.shape)

# =====
#
#   # END YOUR CODE HERE
#   =====
#

    return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

```

```

for i in np.arange(num_checks):
    ix = tuple([np.random.randint(m) for m in self.W.shape])

    oldval = self.W[ix]
    self.W[ix] = oldval + h # increment by h
    fxph = self.loss(X, y)
    self.W[ix] = oldval - h # decrement by h
    fxmh = self.loss(X,y) # evaluate f(x - h)
    self.W[ix] = oldval # reset

    grad_numerical = (fxph - fxmh) / (2 * h)
    grad_analytic = your_grad[ix]
    rel_error = abs(grad_numerical - grad_analytic) /
(abs(grad_numerical) + abs(grad_analytic))
    print('numerical: %f analytic: %f, relative error: %e' %
(grad_numerical, grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and outputs as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # =====
    #
    # YOUR CODE HERE:
    # Calculate the softmax loss and gradient WITHOUT any for loops.
    # =====
    #
    e_x = np.exp(X.dot(self.W.T)) # exp(wT * X) <-- exp of score of
each class (numerator)
    ex_sum = e_x/np.sum(e_x, axis=1, keepdims=True) # ex/sum ex (so
this adds in the denominator) output is (10,)
    # print("ex_sum shape", ex_sum.shape)
    log_sc = np.log(ex_sum) # take the log likelihood

    log_sum = np.sum(log_sc[np.arange(X.shape[0]),y])
    loss = -log_sum
    loss = loss/(X.shape[0]) # normalize the loss
    # print("fast loss", loss)

    grad = (ex_sum).T.dot(X) # this is the default, otherwise if i ==
j then subtract X[i]
    # HOW TO DO MASKING?? for loss when i = j? use np arrange?
    mask_zero=np.zeros_like(ex_sum)
    mask_zero[np.arange(X.shape[0]),y]=1
    mask = ex_sum-mask_zero # mask 500, 10

```

```

    # print("mask:", mask.shape)
    masked_grad = mask.T.dot(X) # X 500, 3072
    # grad -= masked_grad
    grad = masked_grad/(X.shape[0]) # normalize the grad # 10, 3073
    # print("grad***", grad.shape)

    # =====
#
    # END YOUR CODE HERE
    # =====
#

    return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there
are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i]
= c
        means that X[i] has label 0 <= c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each
step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training
iteration.
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where
K is number of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes
the weights of self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

```

```

#
===== #
# YOUR CODE HERE:
#   Sample batch_size elements from the training data for use in
#   gradient descent. After sampling,
#   - X_batch should have shape: (batch_size, dim)
#   - y_batch should have shape: (batch_size,)
#   The indices should be randomly generated to reduce
correlations
#   in the dataset. Use np.random.choice. It's okay to sample
with
#   replacement.
#
===== #
# X_batch = [batch_size][dim]
# y_batch = [batch_size]

# rand_list = np.random.choice(X.shape[0], batch_size)
# for i in rand_list:
#   X_batch.append(X[i])
# rand_list = np.random.choice(y.shape[0], batch_size)
# for i in rand_list:
#   y_batch.append(y[i])

b_samples = []
b_labels = []

for i in range(batch_size):
    index = np.random.choice(num_train)
    b_samples.append(X[index])
    b_labels.append(y[index])

X_batch = np.array(b_samples)
y_batch = np.array(b_labels)

#
===== #
# END YOUR CODE HERE
#
===== #

# evaluate loss and gradient
loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
loss_history.append(loss)

#
===== #
# YOUR CODE HERE:
#   Update the parameters, self.W, with a gradient step

```

```

#
===== #
    self.W = self.W - grad * learning_rate

#
===== #
    # END YOUR CODE HERE
#
===== #

    if verbose and it % 100 == 0:
        print('iteration {} / {}: loss {}'.format(it, num_iters,
loss))

    return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional
point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-
dimensional
    array of length N, and each element is an integer giving the
predicted
    class.
    """
    y_pred = np.zeros(X.shape[0])

# =====
#
# YOUR CODE HERE:
#   Predict the labels given the training data.
# =====
#
# print("ypred shape", y_pred.shape)
# X = X[:3073]
p = (self.W @ X.T).T
# print("pshape", p.shape)
# print("xshape", X.shape)
# print("wshape", self.W.shape)

for i in range(y_pred.shape[0]):
    y_pred[i] = np.argmax(p[i])
# =====
#
# END YOUR CODE HERE
# =====

```

```
#
```

```
    return y_pred
```