

```

import numpy as np

from .layers import *
from .layer_utils import *

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network with ReLU nonlinearity
    and
    softmax loss that uses a modular layer design. We assume an input
    dimension
    of D, a hidden dimension of H, and perform classification over C
    classes.

    The architecture should be affine - relu - affine - softmax.

    Note that this class does not implement gradient descent; instead,
    it
    will interact with a separate Solver object that is responsible for
    running
    optimization.

    The learnable parameters of the model are stored in the dictionary
    self.params that maps parameter names to numpy arrays.
    """

    def __init__(self, input_dim=3*32*32, hidden_dims=100,
num_classes=10,
                dropout=0, weight_scale=1e-3, reg=0.0):
        """
        Initialize a new network.

        Inputs:
        - input_dim: An integer giving the size of the input
        - hidden_dims: An integer giving the size of the hidden layer
        - num_classes: An integer giving the number of classes to classify
        - dropout: Scalar between 0 and 1 giving dropout strength.
        - weight_scale: Scalar giving the standard deviation for random
        initialization of the weights.
        - reg: Scalar giving L2 regularization strength.
        """
        self.params = {}
        self.reg = reg

        # =====
#
        # YOUR CODE HERE:
        # Initialize W1, W2, b1, and b2. Store these as
self.params['W1'],

```

```

        # self.params['W2'], self.params['b1'] and self.params['b2'].
The
        # biases are initialized to zero and the weights are initialized
        # so that each parameter has mean 0 and standard deviation
weight_scale.
        # The dimensions of W1 should be (input_dim, hidden_dim) and the
        # dimensions of W2 should be (hidden_dims, num_classes)
        # =====
#

        self.params = {}
        self.params['W1'] = weight_scale * np.random.randn(input_dim,
hidden_dims)
        self.params['b1'] = np.zeros(hidden_dims)
        self.params['W2'] = weight_scale * np.random.randn(hidden_dims,
num_classes)
        self.params['b2'] = np.zeros(num_classes)

        # =====
#
        # END YOUR CODE HERE
        # =====
#

def loss(self, X, y=None):
    """
    Compute loss and gradient for a minibatch of data.

    Inputs:
    - X: Array of input data of shape (N, d_1, ..., d_k)
    - y: Array of labels, of shape (N,). y[i] gives the label for
X[i].

    Returns:
    If y is None, then run a test-time forward pass of the model and
return:
    - scores: Array of shape (N, C) giving classification scores,
where
        scores[i, c] is the classification score for X[i] and class c.

    If y is not None, then run a training-time forward and backward
pass and
    return a tuple of:
    - loss: Scalar value giving the loss
    - grads: Dictionary with the same keys as self.params, mapping
parameter
        names to gradients of the loss with respect to those parameters.
    """
    scores = None

```

```

# =====
#
# YOUR CODE HERE:
# Implement the forward pass of the two-layer neural network.
Store
# the class scores as the variable 'scores'. Be sure to use the
layers
# you prior implemented.
# =====
#
# Unpack variables from the params dictionary
W1, b1 = self.params['W1'], self.params['b1']
W2, b2 = self.params['W2'], self.params['b2']

h1 = affine_forward(X, W1, b1)
relu_1 = relu_forward(h1[0])
h2 = affine_forward(relu_1[0], W2, b2)
scores = h2[0]
# =====
#
# END YOUR CODE HERE
# =====
#

# If y is None then we are in test mode so just return scores
if y is None:
    return scores

loss, grads = 0, {}
# =====
#
# YOUR CODE HERE:
# Implement the backward pass of the two-layer neural net.
Store
# the loss as the variable 'loss' and store the gradients in the
# 'grads' dictionary. For the grads dictionary, grads['W1']
holds
# the gradient for W1, grads['b1'] holds the gradient for b1,
etc.
# i.e., grads[k] holds the gradient for self.params[k].
#
# Add L2 regularization, where there is an added cost
0.5*self.reg*W^2
# for each W. Be sure to include the 0.5 multiplying factor to
# match our implementation.
#
# And be sure to use the layers you prior implemented.
# =====

```

```

#
    # grads['W2'], grads['b2'], grads['W1'], grads['b1'] = None, None,
None, None
    # W1, b1 = self.params['W1'], self.params['b1']
    # W2, b2 = self.params['W2'], self.params['b2']

    softmax = softmax_loss(h2[0], y)
    loss = softmax[0] + 0.5 * self.reg * (np.sum(W1 ** 2) + np.sum(W2
** 2))

    # h1_grad, h2_grad, relu_grad = None, None, None

    # since w2 has problems need to check if h2[0] is correct shape
but i think it is?
    # need to draw out and check dims
    # print("h2[0] shape", h2[0].shape)
    # print("w2 shape", W2.shape)
    # print("b2 shape", b2.shape)
    # print("W1 shape", W1.shape)
    # print("b1 shape", b1.shape)
    # print("relu shape", relu_1[0].shape)
    # print("x shape", X.shape)
    (h2_grad, grads['W2'], grads['b2']) = affine_backward(softmax[1],
(relu_1[0], W2, b2))

    relu_grad = relu_backward(h2_grad, h1[0])

    # print("W2 shape", W2.shape)
    # print("grads['W2'] shape", grads['W2'].shape)

    (h1_grad, grads['W1'], grads['b1']) = affine_backward(relu_grad,
(X, W1, b1))

    # print("grads w2", grads['W2'].shape)
    # print("w2 shape", W2.shape)

    # add regularization
    # print("w1 shape", W1.shape)
    # print("w1 grad shape", grads['W1'].shape)

    grads['W2'] += self.reg*W2
    grads['W1'] += self.reg*W1

# =====
#
# END YOUR CODE HERE
# =====

```

```
#
```

```
    return loss, grads
```

```
class FullyConnectedNet(object):
```

```
    """
```

```
    A fully-connected neural network with an arbitrary number of hidden layers,
```

```
    ReLU nonlinearities, and a softmax loss function. This will also implement
```

```
    dropout and batch normalization as options. For a network with L layers,
```

```
    the architecture will be
```

```
    {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
```

```
    where batch normalization and dropout are optional, and the {...} block is
```

```
    repeated L - 1 times.
```

```
    Similar to the TwoLayerNet above, learnable parameters are stored in the
```

```
    self.params dictionary and will be learned using the Solver class.
```

```
    """
```

```
    def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10, dropout=0, use_batchnorm=False, reg=0.0, weight_scale=1e-2, dtype=np.float32, seed=None):
```

```
        """
```

```
        Initialize a new FullyConnectedNet.
```

```
        Inputs:
```

```
        - hidden_dims: A list of integers giving the size of each hidden layer.
```

```
        - input_dim: An integer giving the size of the input.
```

```
        - num_classes: An integer giving the number of classes to classify.
```

```
        - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
```

```
            the network should not use dropout at all.
```

```
        - use_batchnorm: Whether or not the network should use batch normalization.
```

```
        - reg: Scalar giving L2 regularization strength.
```

```
        - weight_scale: Scalar giving the standard deviation for random initialization of the weights.
```

```
        - dtype: A numpy datatype object; all computations will be performed using
```

```
            this datatype. float32 is faster but less accurate, so you
```

```

should use
    float64 for numeric gradient checking.
    - seed: If not None, then pass this random seed to the dropout
layers. This
    will make the dropout layers deterministic so we can gradient
check the
    model.
    """
    self.use_batchnorm = use_batchnorm
    self.use_dropout = dropout > 0
    self.reg = reg
    self.num_layers = 1 + len(hidden_dims)
    self.dtype = dtype
    self.params = {}

    # =====
#
# YOUR CODE HERE:
# Initialize all parameters of the network in the self.params
dictionary.
# The weights and biases of layer 1 are W1 and b1; and in
general the
# weights and biases of layer i are Wi and bi. The
# biases are initialized to zero and the weights are initialized
# so that each parameter has mean 0 and standard deviation
weight_scale.
# =====
#

    self.params['W1'] = weight_scale * np.random.randn(input_dim,
hidden_dims[0])
    self.params['b1'] = np.zeros(hidden_dims[0])

    for i in range(1, self.num_layers - 1):
        w = str("W" + str(i + 1))
        b = str("b" + str(i + 1))

        self.params[w] = weight_scale *
np.random.randn(hidden_dims[i-1], hidden_dims[i])
        self.params[b] = np.zeros(hidden_dims[i])

    w = str("W" + str(self.num_layers))
    b = str("b" + str(self.num_layers))
    self.params[w] = weight_scale * np.random.randn(hidden_dims[-1],
num_classes)
    self.params[b] = np.zeros(num_classes)

    # for key in self.params:
    #     print(key, self.params[key].shape)

```

```

# =====
#
# END YOUR CODE HERE
# =====
#

# When using dropout we need to pass a dropout_param dictionary to
each
# dropout layer so that the layer knows the dropout probability
and the mode
# (train / test). You can pass the same dropout_param to each
dropout layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
    if seed is not None:
        self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running means
and
# variances, so we need to pass a special bn_param object to each
batch
# normalization layer. You should pass self.bn_params[0] to the
forward pass
# of the first batch normalization layer, self.bn_params[1] to the
forward
# pass of the second batch normalization layer, etc.
self.bn_params = []
if self.use_batchnorm:
    self.bn_params = [{'mode': 'train'} for i in
np.arange(self.num_layers - 1)]

# Cast all parameters to the correct datatype
for k, v in self.params.items():
    self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param since
they
    # behave differently during training and testing.
    if self.dropout_param is not None:

```

```

        self.dropout_param['mode'] = mode
    if self.use_batchnorm:
        for bn_param in self.bn_params:
            bn_param[mode] = mode

    scores = None

    # =====
#
# YOUR CODE HERE:
#   Implement the forward pass of the FC net and store the output
#   scores as the variable "scores".
# =====
#

# forward prop
Hs = [X]
Zs = [X]
Ws = []
bs = []

W = self.params['W1']
b = self.params['b1']
aff_fwd = affine_forward(X, W, b)
Z = aff_fwd[0]

for i in range(1, self.num_layers):
    relu_h = relu_forward(Z)
    H = relu_h[0]
    Hs.append(H)
    Ws.append(W)
    Zs.append(Z)
    bs.append(b)

    H = Hs[-1]
    W = self.params[str('W' + str(i+1))]
    b = self.params[str('b' + str(i+1))]

    aff_fwd = affine_forward(H, W, b)
    Z = aff_fwd[0]
    scores = Z

    Zs.append(Z)
    Ws.append(W)
    bs.append(b)
    # print("done with fwd pass")

# =====
#

```



```

# END YOUR CODE HERE
# =====
#
# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# =====
#
# YOUR CODE HERE:
# Implement the backwards pass of the FC net and store the
gradients
# in the grads dict, so that grads[k] is the gradient of
self.params[k]
# Be sure your L2 regularization includes a 0.5 factor.
# =====
#

softmax = softmax_loss(Zs[-1], y)

# regularization
agg_sum = 0
for W in Ws:
    agg_sum += np.sum(W ** 2)

loss = softmax[0] + 0.5 * self.reg * agg_sum

# backprop
loss_grads = [softmax[1]]

for i in range(self.num_layers, 1, -1):
    loss_grad = loss_grads[-1]
    (h_grad, grads[str('W' + str(i))], grads[str('b' + str(i))]) =
affine_backward(loss_grad, (Hs[i-1], Ws[i-1], bs[i-1]))
    relu_grad = relu_backward(h_grad, Zs[i-1])
    loss_grads.append(relu_grad)

loss_grad = loss_grads[-1]
(h_grad, grads['W1'], grads['b1']) = affine_backward(loss_grad,
(X, Ws[0], bs[0]))

# regularization
for i in range(self.num_layers):
    grads[str('W' + str(i+1))] += self.reg * self.params[str('W' +
str(i+1))]

# for key in grads.keys():

```

```

# print(key, grads[key].shape)

# loss_grad = loss_grads[-1]
# (h_grad, grads['W3'], grads['b3']) = affine_backward(loss_grad,
(Zs[-1], Ws[-1], bs[-1]))
# print("loss grad shape", loss_grad[0].shape)
# print("h grad shape", h_grad.shape)

# # print("w grad size", w_grad.shape)

# for i in range(self.num_layers - 1, 0, -1):
#     print(i)
#     loss_grad = loss_grads[-1]
#     print("h_grad size", h_grad.shape)
#     print("Hs[i] shape", Hs[i-1].shape)
#     relu_grad = relu_backward(h_grad, Hs[i-1])
#     print("relu grad shape", relu_grad.shape)
#     print("h shape", Hs[i-1].shape)
#     print("w shape", Ws[i-1].shape)
#     print("b shape", bs[i-1].shape)
#     (h_grad, w_grad, b_grad) = affine_backward(relu_grad,
(Hs[i-1], Ws[i-1], bs[i-1]))
#     # print("new h shape", h_grad.shape)
#     loss_grads.append(h_grad)
#     print(str('W' + str(i)))
#     grads[str('W' + str(i))] = w_grad
#     grads[str('b' + str(i))] = b_grad

# regularization

# print("i", i)
# print(str('W'+str(i+1)))
# # need to figure this part out
# print(grads['W1'])
# #self.reg*self.params['W1']
# print("i got here 2")

# j = 0
# for W in Ws:
#     ≈
#     j = j+1

```

```

    # (h2_grad, grads['W2'], grads['b2']) =
    affine_backward(softmax[1], (relu_1[0], W2, b2))

    # relu_grad = relu_backward(h2_grad, h1[0])

    # # print("W2 shape", W2.shape)
    # # print("grads['W2'] shape", grads['W2'].shape)

    # (h1_grad, grads['W1'], grads['b1']) = affine_backward(relu_grad,
    (X, W1, b1))

    # # print("grads w2", grads['W2'].shape)
    # # print("w2 shape", W2.shape)

    # # add regularization
    # # print("w1 shape", W1.shape)
    # # print("w1 grad shape", grads['W1'].shape)

    # grads['W2'] += self.reg*W2
    # grads['W1'] += self.reg*W1

    # =====
# # END YOUR CODE HERE
# =====
#
    return loss, grads

```