

```
import numpy as np
import pdb
```

```
def affine_forward(x, w, b):
    """
```

Computes the forward pass for an affine (fully-connected) layer.

The input x has shape (N, d_1, \dots, d_k) and contains a minibatch of N examples, where each example $x[i]$ has shape (d_1, \dots, d_k) . We will reshape each input into a vector of dimension $D = d_1 * \dots * d_k$, and then transform it to an output vector of dimension M .

Inputs:

- x : A numpy array containing input data, of shape (N, d_1, \dots, d_k)
- w : A numpy array of weights, of shape (D, M)
- b : A numpy array of biases, of shape $(M,)$

Returns a tuple of:

- out : output, of shape (N, M)
- $cache$: (x, w, b)

```
    """
```

```
    # ===== #
    # YOUR CODE HERE:
    # Calculate the output of the forward pass. Notice the dimensions
    # of  $w$  are  $D \times M$ , which is the transpose of what we did in earlier
    # assignments.
    # ===== #
```

```
    # print("w", w.shape)
    # print("xr", x.shape)
    # print("b", b.shape)
    # print("x shape", x.shape)
    xr = x.reshape(x.shape[0], -1)
    # print("x reshape", xr.shape)
```

```
    out = xr.dot(w) + b
    # print("wxr", xr.dot(w).shape) # check this shape to see if h2
    shape is even correct???
```

```
    # ===== #
    # END YOUR CODE HERE
    # ===== #
```

```

cache = (x, w, b)
return out, cache

```

```

def affine_backward(dout, cache):

```

```

    """

```

```

    Computes the backward pass for an affine layer.

```

```

    Inputs:

```

- dout: Upstream derivative, of shape (N, M)
- cache: Tuple of:
 - x: Input data, of shape (N, d_1, ... d_k)
 - w: Weights, of shape (D, M)

```

    Returns a tuple of:

```

- dx: Gradient with respect to x, of shape (N, d1, ..., d_k)
- dw: Gradient with respect to w, of shape (D, M)
- db: Gradient with respect to b, of shape (M,)

```

    """

```

```

    x, w, b = cache

```

```

    dx, dw, db = None, None, None

```

```

    # ===== #

```

```

    # YOUR CODE HERE:

```

```

    # Calculate the gradients for the backward pass.

```

```

    # ===== #

```

```

    # dout is N x M

```

```

    # dx should be N x d1 x ... x dk; it relates to dout through
multiplication with w, which is D x M

```

```

    # dw should be D x M; it relates to dout through multiplication with
x, which is N x D after reshaping

```

```

    # db should be M; it is just the sum over dout examples

```

```

    # dout: 5 x 3

```

```

    # w: 3 x 10

```

```

    # x: 10 x 5 which is M x N

```

```

    # print("x shape", x.shape)

```

```

    # n = 5

```

```

    # m =

```

```

    # d = 3

```

```

    # m = THERE IS SMTH WRONG HERE!!! WITH THE NUMBERS

```

```

    # print("dout", dout.shape)

```

```

    # print("w", w.shape)

```

```

    dx = np.dot(dout, w.T) # (N, M) x (M, D) = (N, D) should be D,M

```

```

    # print("dout shape", dout.shape)

```

```

    # print("w.T shape", w.T.shape)

```

```

    # print("x shape", x.shape)

```

```

    # dx = dx.reshape(x.shape)

    # dw = np.dot(x.reshape(x.shape[0], -1).T, dout) # (N, M) x (1 x N)
    should be (D, M) original
    xr = x.reshape(x.shape[0], -1)
    dw = np.dot(xr.T, dout) #not sure if this works
    db = np.sum(dout, axis=0) # M i think this one is fine
    # print("x shape", x.shape)
    # print("dx shape", dx.shape, w.shape[0], dout.shape[1])
    # print("dw shape", dw.shape, dout.shape[0], w.shape[0])
    # print("db shape", db.shape, dout.shape[1])

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return dx, dw, db

def relu_forward(x):
    """
    Computes the forward pass for a layer of rectified linear units
    (ReLU).

    Input:
    - x: Inputs, of any shape

    Returns a tuple of:
    - out: Output, of the same shape as x
    - cache: x
    """
    # ===== #
    # YOUR CODE HERE:
    # Implement the ReLU forward pass.
    # ===== #

    f = lambda x: x * (x > 0)
    out = f(x)
    # print("out", out.shape)
    # print("x", x.shape)
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    cache = x
    return out, cache

def relu_backward(dout, cache):
    """

```

Computes the backward pass for a layer of rectified linear units (ReLU).

Input:

- dout: Upstream derivatives, of any shape
- cache: Input x, of same shape as dout

Returns:

- dx: Gradient with respect to x

"""

x = cache

=====

YOUR CODE HERE:

Implement the ReLU backward pass

=====

ReLU directs linearly to those > 0

xr = x.reshape(x.shape[0], -1)

dx = dout * (x > 0) # hadamard product is element wise operation

=====

END YOUR CODE HERE

=====

return dx

def softmax_loss(x, y):

"""

Computes the loss and gradient for softmax classification.

Inputs:

- x: Input data, of shape (N, C) where x[i, j] is the score for the jth class

for the ith input.

- y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and

0 <= y[i] < C

Returns a tuple of:

- loss: Scalar giving the loss

- dx: Gradient of the loss with respect to x

"""

probs = np.exp(x - np.max(x, axis=1, keepdims=True))

probs /= np.sum(probs, axis=1, keepdims=True)

N = x.shape[0]

loss = -np.sum(np.log(probs[np.arange(N), y])) / N

dx = probs.copy()

```
dx[np.arange(N), y] -= 1  
dx /= N  
return loss, dx
```