

```
import numpy as np
import matplotlib.pyplot as plt
from .layer_utils import *
```

```
class TwoLayerNet(object):
    """
```

A two-layer fully-connected neural network. The net has an input dimension of D , a hidden layer dimension of H , and performs classification over C classes.

We train the network with a softmax loss function and L2 regularization on the weight matrices. The network uses a ReLU nonlinearity after the first fully connected layer.

In other words, the network has the following architecture:

input – fully connected layer – ReLU – fully connected layer – softmax

The outputs of the second fully-connected layer are the scores for each class.

```
    """
    def __init__(self, input_size, hidden_size, output_size, std=1e-4):
        """
```

Initialize the model. Weights are initialized to small random values and biases are initialized to zero. Weights and biases are stored in the variable `self.params`, which is a dictionary with the following keys:

- W1: First layer weights; has shape (H, D)
- b1: First layer biases; has shape $(H,)$
- W2: Second layer weights; has shape (C, H)
- b2: Second layer biases; has shape $(C,)$

Inputs:

- `input_size`: The dimension D of the input data.
- `hidden_size`: The number of neurons H in the hidden layer.
- `output_size`: The number of classes C .

```
    """
    self.params = {}
    self.params['W1'] = std * np.random.randn(hidden_size, input_size)
    self.params['b1'] = np.zeros(hidden_size)
    self.params['W2'] = std * np.random.randn(output_size,
hidden_size)
```

```

self.params['b2'] = np.zeros(output_size)

def loss(self, X, y=None, reg=0.0):
    """
    Compute the loss and gradients for a two layer fully connected
    neural
    network.

    Inputs:
    - X: Input data of shape (N, D). Each X[i] is a training sample.
    - y: Vector of training labels. y[i] is the label for X[i], and
    each y[i] is
        an integer in the range  $0 \leq y[i] < C$ . This parameter is
    optional; if it
        is not passed then we only return scores, and if it is passed
    then we
        instead return the loss and gradients.
    - reg: Regularization strength.

    Returns:
    If y is None, return a matrix scores of shape (N, C) where
    scores[i, c] is
        the score for class c on input X[i].

    If y is not None, instead return a tuple of:
    - loss: Loss (data loss and regularization loss) for this batch of
    training
        samples.
    - grads: Dictionary mapping parameter names to gradients of those
    parameters
        with respect to the loss function; has the same keys as
    self.params.
    """
    # Unpack variables from the params dictionary
    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']
    N, D = X.shape

    # Compute the forward pass
    scores = None

    # =====
#
# YOUR CODE HERE:
# Calculate the output scores of the neural network. The result
# should be (N, C). As stated in the description for this class,
# there should not be a ReLU layer after the second FC layer.
# The output of the second FC layer is the output scores. Do not
# use a for loop in your implementation.

```

```

# =====
#
f = lambda x: x * (x > 0) # relu

h1 = W1.dot(X.T) + b1.reshape(b1.shape[0], 1) # (H, D) x (D, N) +
(H, ) --> (H, N)
relu_1 = f(h1)
h2 = W2.dot(relu_1) + b2.reshape(b2.shape[0], 1) # (C, H) x (H, N)
--> (C, N)
scores = h2.T # (C, N) --> (N, C)

# =====
#
# END YOUR CODE HERE
# =====
#

# If the targets are not given then jump out, we're done
if y is None:
    return scores

# Compute the loss
loss = None

# =====
#
# YOUR CODE HERE:
# Calculate the loss of the neural network. This includes the
# softmax loss and the L2 regularization for W1 and W2. Store
the
# total loss in the variable loss. Multiply the regularization
# loss by 0.5 (in addition to the factor reg).
# =====
#
#
# scores is num_examples by num_classes
s_loss = softmax_loss(h2.T, y)[0]
L2_reg = 0.5 * reg * (np.sum(W1 ** 2) + np.sum(W2 ** 2))
loss = s_loss + L2_reg
# =====
#
# END YOUR CODE HERE
# =====
#

grads = {}

```

```

# =====
#
# YOUR CODE HERE:
# Implement the backward pass. Compute the derivatives of the
# weights and the biases. Store the results in the grads
# dictionary. e.g., grads['W1'] should store the gradient for
# W1, and be of the same size as W1.
# =====
#

d1 = softmax_loss(h2.T, y)[1]

# print("w1 dim", W1.shape)
# print("w2 dim", W2.shape)
# print("relu_1 dim", relu_1.shape)
# print("d1 dim", d1.shape)

grads['b2'] = sum(d1)
d2 = d1 @ W2
# print("i got here")
grads['W2'] = d1.T @ relu_1.T + reg * W2

d3 = (h1 > 0) * d2.T
grads['b1'] = sum(d3.T)

# print("w1 dim", W1.shape)
# print("x dim", X.shape)
# print("d3 dim", d3.shape)
grads['W1'] = d3 @ X + reg * W1

# =====
#
# END YOUR CODE HERE
# =====
#

return loss, grads

def train(self, X, y, X_val, y_val,
          learning_rate=1e-3, learning_rate_decay=0.95,
          reg=1e-5, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this neural network using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) giving training data.

```

```

- y: A numpy array of shape (N,) giving training labels; y[i] = c
means that
    X[i] has label c, where  $0 \leq c < C$ .
- X_val: A numpy array of shape (N_val, D) giving validation data.
- y_val: A numpy array of shape (N_val,) giving validation labels.
- learning_rate: Scalar giving learning rate for optimization.
- learning_rate_decay: Scalar giving factor used to decay the
learning rate
    after each epoch.
- reg: Scalar giving regularization strength.
- num_iters: Number of steps to take when optimizing.
- batch_size: Number of training examples to use per step.
- verbose: boolean; if true print progress during optimization.
"""
num_train = X.shape[0]
iterations_per_epoch = max(num_train / batch_size, 1)

# Use SGD to optimize the parameters in self.model
loss_history = []
train_acc_history = []
val_acc_history = []

for it in np.arange(num_iters):
    X_batch = None
    y_batch = None

    #
    ===== #
    # YOUR CODE HERE:
    #   Create a minibatch by sampling batch_size samples randomly.
    #
    ===== #
    b_samples = []
    b_labels = []

    for i in range(batch_size):
        index = np.random.choice(num_train)
        b_samples.append(X[index])
        b_labels.append(y[index])

    X_batch = np.array(b_samples)
    y_batch = np.array(b_labels)

    #
    ===== #
    # END YOUR CODE HERE
    #
    ===== #

    # Compute loss and gradients using the current minibatch

```

```

        loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
        loss_history.append(loss)

#
===== #
# YOUR CODE HERE:
#   Perform a gradient descent step using the minibatch to
update
#   all parameters (i.e., W1, W2, b1, and b2).
#
===== #

# Unpack variables from the params dictionary
W1, b1 = self.params['W1'], self.params['b1']
W2, b2 = self.params['W2'], self.params['b2']

# update
self.params['W2'] = W2 - grads['W2'] * learning_rate
self.params['W1'] = W1 - grads['W1'] * learning_rate
self.params['b2'] = b2 - grads['b2'] * learning_rate
self.params['b1'] = b1 - grads['b1'] * learning_rate

#
===== #
# END YOUR CODE HERE
#
===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters,
loss))

# Every epoch, check train and val accuracy and decay learning
rate.
if it % iterations_per_epoch == 0:
    # Check accuracy
    train_acc = (self.predict(X_batch) == y_batch).mean()
    val_acc = (self.predict(X_val) == y_val).mean()
    train_acc_history.append(train_acc)
    val_acc_history.append(val_acc)

    # Decay learning rate
    learning_rate *= learning_rate_decay

return {
    'loss_history': loss_history,
    'train_acc_history': train_acc_history,
    'val_acc_history': val_acc_history,
}

```

```

def predict(self, X):
    """
    Use the trained weights of this two-layer network to predict
    labels for
    data points. For each data point we predict scores for each of the
    C
    classes, and assign each data point to the class with the highest
    score.

    Inputs:
    - X: A numpy array of shape (N, D) giving N D-dimensional data
    points to
    classify.

    Returns:
    - y_pred: A numpy array of shape (N,) giving predicted labels for
    each of
    the elements of X. For all i, y_pred[i] = c means that X[i] is
    predicted
    to have class c, where 0 <= c < C.
    """
    y_pred = None

    # =====
    #
    # YOUR CODE HERE:
    # Predict the class given the input data.
    # =====
    #
    # p = (self.W @ X.T).T

    y_pred = np.zeros(X.shape[0])

    # Unpack variables from the params dictionary
    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']

    f = lambda x: x * (x > 0) # relu

    h1 = W1.dot(X.T) + b1.reshape(b1.shape[0], 1) # (H, D) x (D, N) +
    (H, ) --> (H, N)
    relu_1 = f(h1)
    h2 = W2.dot(relu_1) + b2.reshape(b2.shape[0], 1) # (C, H) x (H, N)
    --> (C, N)
    p = h2.T # (C, N) --> (N, C)

    for i in range(y_pred.shape[0]):

```

```
        y_pred[i] = np.argmax(p[i])

# =====
#
# END YOUR CODE HERE
# =====
#

return y_pred
```