

```

import numpy as np
import pdb

from .layers import *
from .layer_utils import *

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network with ReLU nonlinearity
    and softmax loss that uses a modular layer design. We assume an input
    dimension of D, a hidden dimension of H, and perform classification over C
    classes.

    The architecture should be affine - relu - affine - softmax.

    Note that this class does not implement gradient descent; instead,
    it will interact with a separate Solver object that is responsible
    for running optimization.

    The learnable parameters of the model are stored in the dictionary
    self.params that maps parameter names to numpy arrays.
    """

    def __init__(self, input_dim=3*32*32, hidden_dims=100,
num_classes=10, dropout=1, weight_scale=1e-3, reg=0.0):
        """
        Initialize a new network.

        Inputs:
        - input_dim: An integer giving the size of the input
        - hidden_dims: An integer giving the size of the hidden layer
        - num_classes: An integer giving the number of classes to
classify
        - dropout: Scalar between 0 and 1 giving dropout strength.
        - weight_scale: Scalar giving the standard deviation for
random
        initialization of the weights.
        - reg: Scalar giving L2 regularization strength.
        """
        self.params = {}
        self.reg = reg

        #
===== #

```

```

        # YOUR CODE HERE:
        # Initialize W1, W2, b1, and b2. Store these as
self.params['W1'],
        # self.params['W2'], self.params['b1'] and
self.params['b2']. The
        # biases are initialized to zero and the weights are
initialized
        # so that each parameter has mean 0 and standard deviation
weight_scale.
        # The dimensions of W1 should be (input_dim, hidden_dim) and
the
        # dimensions of W2 should be (hidden_dims, num_classes)
        #
===== #

        self.params['W1'] = weight_scale * np.random.randn(input_dim,
hidden_dims)
        self.params['b1'] = np.zeros(hidden_dims)
        self.params['W2'] = weight_scale *
np.random.randn(hidden_dims, num_classes)
        self.params['b2'] = np.zeros(num_classes)

        #
===== #
        # END YOUR CODE HERE
        #
===== #

def loss(self, X, y=None):
    """
    Compute loss and gradient for a minibatch of data.

    Inputs:
    - X: Array of input data of shape (N, d_1, ..., d_k)
    - y: Array of labels, of shape (N,). y[i] gives the label for
X[i].

    Returns:
    If y is None, then run a test-time forward pass of the model
and return:
    - scores: Array of shape (N, C) giving classification scores,
where
        scores[i, c] is the classification score for X[i] and class
c.

    If y is not None, then run a training-time forward and
backward pass and
    return a tuple of:
    - loss: Scalar value giving the loss
    - grads: Dictionary with the same keys as self.params, mapping

```

```

parameter
    names to gradients of the loss with respect to those
parameters.
    """
    scores = None

    #
===== #
    # YOUR CODE HERE:
    # Implement the forward pass of the two-layer neural
network. Store
    # the class scores as the variable 'scores'. Be sure to use
the layers
    # you prior implemented.
    #
===== #

    # Unpack variables from the params dictionary
    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']

    h1 = affine_forward(X, W1, b1)
    relu_1 = relu_forward(h1[0])
    h2 = affine_forward(relu_1[0], W2, b2)
    scores = h2[0]

    #
===== #
    # END YOUR CODE HERE
    #
===== #

    # If y is None then we are in test mode so just return scores
    if y is None:
        return scores

    loss, grads = 0, {}
    #
===== #
    # YOUR CODE HERE:
    # Implement the backward pass of the two-layer neural net.
Store
    # the loss as the variable 'loss' and store the gradients in
the
    # 'grads' dictionary. For the grads dictionary, grads['W1']
holds
    # the gradient for W1, grads['b1'] holds the gradient for
b1, etc.
    # i.e., grads[k] holds the gradient for self.params[k].

```

```

        #
        # Add L2 regularization, where there is an added cost
0.5*self.reg*W^2
        # for each W. Be sure to include the 0.5 multiplying factor
to
        # match our implementation.
        #
        # And be sure to use the layers you prior implemented.
        #
===== #

        softmax = softmax_loss(h2[0], y)
        loss = softmax[0] + 0.5 * self.reg * (np.sum(W1 ** 2) +
np.sum(W2 ** 2))

        (h2_grad, grads['W2'], grads['b2']) =
affine_backward(softmax[1], (relu_1[0], W2, b2))

        relu_grad = relu_backward(h2_grad, h1[0])

        (h1_grad, grads['W1'], grads['b1']) =
affine_backward(relu_grad, (X, W1, b1))

        grads['W2'] += self.reg*W2
        grads['W1'] += self.reg*W1

        #
===== #
        # END YOUR CODE HERE
        #
===== #

    return loss, grads

```

```

class FullyConnectedNet(object):
    """
    A fully-connected neural network with an arbitrary number of
hidden layers,
    ReLU nonlinearities, and a softmax loss function. This will also
implement
    dropout and batch normalization as options. For a network with L
layers,
    the architecture will be

    {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine -
softmax

    where batch normalization and dropout are optional, and the {...}
block is

```

repeated  $L - 1$  times.

Similar to the TwoLayerNet above, learnable parameters are stored in the

`self.params` dictionary and will be learned using the Solver class.

```
def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
              dropout=1, use_batchnorm=False, reg=0.0,
              weight_scale=1e-2, dtype=np.float32, seed=None):
    """
```

Initialize a new FullyConnectedNet.

Inputs:

- `hidden_dims`: A list of integers giving the size of each hidden layer.
- `input_dim`: An integer giving the size of the input.
- `num_classes`: An integer giving the number of classes to classify.
- `dropout`: Scalar between 0 and 1 giving dropout strength. If `dropout=1` then the network should not use dropout at all.
- `use_batchnorm`: Whether or not the network should use batch normalization.
- `reg`: Scalar giving L2 regularization strength.
- `weight_scale`: Scalar giving the standard deviation for random initialization of the weights.
- `dtype`: A numpy datatype object; all computations will be performed using this datatype. `float32` is faster but less accurate, so you should use `float64` for numeric gradient checking.
- `seed`: If not None, then pass this random seed to the dropout layers. This will make the dropout layers deterministic so we can gradient check the model.

```
    """
    self.use_batchnorm = use_batchnorm
    self.use_dropout = dropout < 1
    self.reg = reg
    self.num_layers = 1 + len(hidden_dims)
    self.dtype = dtype
    self.params = {}
```

#

===== #

# YOUR CODE HERE:

# Initialize all parameters of the network in the

```

self.params dictionary.
# The weights and biases of layer 1 are W1 and b1; and in
general the
# weights and biases of layer i are Wi and bi. The
# biases are initialized to zero and the weights are
initialized
# so that each parameter has mean 0 and standard deviation
weight_scale.
#
# BATCHNORM: Initialize the gammas of each layer to 1 and
the beta
# parameters to zero. The gamma and beta parameters for
layer 1 should
# be self.params['gamma1'] and self.params['beta1']. For
layer 2, they
# should be gamma2 and beta2, etc. Only use batchnorm if
self.use_batchnorm
# is true and DO NOT do batch normalize the output scores.
#
===== #

    self.params['W1'] = weight_scale * np.random.randn(input_dim,
hidden_dims[0])
    self.params['b1'] = np.zeros(hidden_dims[0])

    for i in range(1, self.num_layers - 1):
        w = str("W" + str(i + 1))
        b = str("b" + str(i + 1))
        self.params[w] = weight_scale *
np.random.randn(hidden_dims[i-1], hidden_dims[i])
        self.params[b] = np.zeros(hidden_dims[i])

    w = str("W" + str(self.num_layers))
    b = str("b" + str(self.num_layers))
    self.params[w] = weight_scale *
np.random.randn(hidden_dims[-1], num_classes)
    self.params[b] = np.zeros(num_classes)

    # batch norm
    if self.use_batchnorm:
        for i in range(self.num_layers - 1):
            self.params['gamma'+str(i+1)] = np.ones(hidden_dims[i])
            self.params['beta'+str(i+1)] = np.zeros(hidden_dims[i])

    #
===== #

```

```

# END YOUR CODE HERE
#
===== #

# When using dropout we need to pass a dropout_param
dictionary to each
# dropout layer so that the layer knows the dropout
probability and the mode
# (train / test). You can pass the same dropout_param to each
dropout layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
if seed is not None:
    self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running
means and
# variances, so we need to pass a special bn_param object to
each batch
# normalization layer. You should pass self.bn_params[0] to
the forward pass
# of the first batch normalization layer, self.bn_params[1] to
the forward
# pass of the second batch normalization layer, etc.
self.bn_params = []
if self.use_batchnorm:
    self.bn_params = [{'mode': 'train'} for i in
np.arange(self.num_layers - 1)]

# Cast all parameters to the correct datatype
for k, v in self.params.items():
    # if type(v) != int:
    self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param
since they
# behave differently during training and testing.
if self.dropout_param is not None:
    self.dropout_param['mode'] = mode

```

```

        if self.use_batchnorm:
            for bn_param in self.bn_params:
                bn_param['mode'] = mode

        scores = None

        #
===== #
        # YOUR CODE HERE:
        #   Implement the forward pass of the FC net and store the
output    #   scores as the variable "scores".
        #
        #   BATCHNORM: If self.use_batchnorm is true, insert a
batchnorm #   between the affine_forward and relu_forward layers. You
layer      #   also write an affine_batchnorm_relu() function in
may        #   layer_utils.py.
        #
        #   DROPOUT: If dropout is non-zero, insert a dropout layer
after      #   every ReLU layer.
        #
===== #

        # try again
        layer_input = X
        relu_cache = []
        affine_cache = []
        bn_cache = []
        drop_cache = []

        for i in range(self.num_layers):
            layer_out, acache = affine_forward(layer_input,
self.params['W'+str(i+1)], self.params['b'+str(i+1)])
            affine_cache.append(acaache)

            # batchnorm/dropout maybe
            if i != self.num_layers - 1:
                if self.use_batchnorm:
                    layer_out, bcache = batchnorm_forward(layer_out,
self.params['gamma'+str(i+1)], self.params['beta'+str(i+1)],
self.bn_params[i])
                    bn_cache.append(bcache)

            # relu
            layer_out, rcache = relu_forward(layer_out)
            relu_cache.append(rcache)

```



```

        # potentially dropout
        if self.use_dropout:
            layer_out, dcache = dropout_forward(layer_out,
self.dropout_param)
            drop_cache.append(dcache)

    layer_input = layer_out

    # last layer scores
    scores = layer_out

    # if self.use_batchnorm:
    #     # use batchnorm
    #     print("using batchnorm")

    #     Hs = [X]
    #     caches = []

    #     print("num layers", self.num_layers)

    #     # { affine - batchnorm - relu } * ( L-1 )
    #     for i in range(self.num_layers - 1):
    #         # print(i)
    #         W = self.params[str('W' + str(i+1))]
    #         b = self.params[str('b' + str(i+1))]
    #         gamma = self.params[str('gamma' + str(i+1))]
    #         beta = self.params[str('beta' + str(i+1))]
    #         H = Hs[-1]
    #         out, cache = affine_batchnorm_relu_forward(H, W, b,
gamma, beta, self.bn_params[i])
    #         # print("w shape", W.shape)
    #         # print("b shape", b.shape)
    #         # print("x shape", H.shape)

    #         Hs.append(out)
    #         caches.append(cache)

    #     # affine - (softmax is in the next part)

    #     # print(i+1)
    #     W = self.params[str('W' + str(i+2))]
    #     b = self.params[str('b' + str(i+2))]
    #     H = Hs[-1]
    #     # print("w shape", W.shape)
    #     # print("b shape", b.shape)
    #     # print("x shape", H.shape)
    #     out, cache = affine_forward(H, W, b)
    #     Hs.append(out)

```

```

#     caches.append(cache)

# # forward prop w/o batch norm
# else:
#     Hs = [X]
#     Zs = [X]
#     Ws = []
#     bs = []

#     W = self.params['W1']
#     b = self.params['b1']
#     aff_fwd = affine_forward(X, W, b)
#     Z = aff_fwd[0]

#     for i in range(1, self.num_layers):
#         relu_h = relu_forward(Z)
#         H = relu_h[0]
#         Hs.append(H)
#         Ws.append(W)
#         Zs.append(Z)
#         bs.append(b)

#         H = Hs[-1]
#         W = self.params[str('W' + str(i+1))]
#         b = self.params[str('b' + str(i+1))]

#         aff_fwd = affine_forward(H, W, b)
#         Z = aff_fwd[0]
#         scores = Z

#     Zs.append(Z)
#     Ws.append(W)
#     bs.append(b)

#
===== #
# END YOUR CODE HERE
#
===== #

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
#

```

```

===== #
    # YOUR CODE HERE:
    #   Implement the backwards pass of the FC net and store the
gradients
    #   in the grads dict, so that grads[k] is the gradient of
self.params[k]
    #   Be sure your L2 regularization includes a 0.5 factor.
    #
    #   BATCHNORM: Incorporate the backward pass of the batchnorm.
    #
    #   DROPOUT: Incorporate the backward pass of dropout.
    #
===== #

    # print('num layers', self.num_layers)

    loss, grad = softmax_loss(scores, y)
    up_grad = grad

    for i in reversed(np.arange(self.num_layers)):
        # print(i)
        # add regularization to loss
        loss += 0.5 * self.reg * np.sum((self.params['W'+str(i+1)] *
self.params['W'+str(i+1)]))

        # backward pass
        if i == 0:
            da, dw, db = affine_backward(up_grad, affine_cache.pop())
        else:
            da, dw, db = affine_backward(up_grad, affine_cache.pop())

            if self.use_dropout:
                da = dropout_backward(da, drop_cache.pop())

            dz = relu_backward(da, relu_cache.pop())
            if self.use_batchnorm:
                dz, dgamma, dbeta = batchnorm_backward(dz,
bn_cache.pop())
                grads['gamma'+str(i)] = dgamma
                grads['beta'+str(i)] = dbeta

            grads['W'+str(i+1)] = dw + self.reg *
self.params['W'+str(i+1)]
            grads['b'+str(i+1)] = db

            up_grad = dz

    # print(grads.keys())

```

```

# print(grads.keys())

# softmax = softmax_loss(scores, y)

# # regularization
# agg_sum = 0
# for i in range(self.num_layers):
#     agg_sum += np.sum(self.params[str('W' + str(i+1))]) ** 2)

# loss = softmax[0] + 0.5 * self.reg * agg_sum

# print('i got here')

# loss_grads = [softmax[1]]
# loss_grad = loss_grads[-1]

# print("num layers", self.num_layers)

# # first do the affine layer
# print(self.num_layers-1)
# dx, grads['W' + str(self.num_layers)], grads['b' +
str(self.num_layers)] = affine_backward(loss_grad,
affine_cache[self.num_layers-1])
# # print("x shape", affine_cache[self.num_layers][0].shape)
# # print("w shape", affine_cache[self.num_layers][1].shape)
# # print("b shape", affine_cache[self.num_layers-1][2].shape)

# # print("dx shape", dx.shape)
# # print("dw shape", grads['W' +
str(self.num_layers-1)].shape)
# # print("db shape", grads['b' +
str(self.num_layers-1)].shape)

# loss_grads.append(dx)

# # then the remaining layers
# for i in range(self.num_layers-1, 0, -1):
#     print(i-1)
#     loss_grad = loss_grads[-1]
#     relu_grad = relu_backward(loss_grad, relu_cache[i-1])

#     if self.use_batchnorm:
#         d_batchnorm, grads['gamma' + str(i)], grads['beta' +
str(i)] = batchnorm_backward(relu_grad, bn_cache[i-1])
#         dx, grads['W' + str(i)], grads['b' + str(i)] =
affine_backward(d_batchnorm, affine_cache[i-1])

#     else:
#         dx, grads['W' + str(i)], grads['b' + str(i)] =

```

```

affine_backward(loss_grad, affine_cache[i-1])

# # print("x shape", affine_cache[i-1][0].shape)
# # print("w shape", affine_cache[i-1][1].shape)
# # print("b shape", affine_cache[i-1][2].shape)

# # print("dx shape", dx.shape)
# # print("dw shape", grads['W' + str(i-1)].shape)
# # print("db shape", grads['b' + str(i-1)].shape)

# loss_grads.append(dx)

# # # back prop
# # if self.use_batchnorm: # batch norm
# #     loss_grads = [softmax[1]]
# #     loss_grad = loss_grads[-1]
# #     cache = caches[-1]

# #     dx, grads['W' + str(self.num_layers-1)], grads['b' +
str(self.num_layers-1)] = affine_backward(loss_grad, cache)
# #     loss_grads.append(dx)

# #     for i in range(self.num_layers-1, 1, -1):
# #         print(i)
# #         loss_grad = loss_grads[-1]
# #         dx, grads[str('W' + str(i))], grads[str('b' + str(i))]]
= affine_batchnorm_relu_backward(loss_grad, caches[i-1])
# #         loss_grads.append(dx)

# # else: # no batchnorm
# #     loss_grads = [softmax[1]]

# #     for i in range(self.num_layers, 1, -1):
# #         loss_grad = loss_grads[-1]
# #         (h_grad, grads[str('W' + str(i))], grads[str('b' +
str(i))]) = affine_backward(loss_grad, (Hs[i-1], Ws[i-1], bs[i-1]))
# #         relu_grad = relu_backward(h_grad, Zs[i-1])
# #         loss_grads.append(relu_grad)

# #     loss_grad = loss_grads[-1]
# #     (h_grad, grads['W1'], grads['b1']) =
affine_backward(loss_grad, (X, Ws[0], bs[0]))

# print(grads.keys())

# for i in range(self.num_layers):
#     print("i: ", i)
#     print("self params shape", self.params[str('W' +

```

```

str(i+1))].shape)

    # for i in range(self.num_layers):
    #     print("i: ", i)
    #     print("w grad shape", grads[str('W' + str(i+1))].shape)

    # # regularization
    # for i in range(self.num_layers):
    #     print("i: ", i)
    #     print("grad shape", grads[str('W' + str(i+1))].shape)
    #     print("w shape", self.params[str('W' + str(i+1))].shape)
    #     grads[str('W' + str(i+1))] += self.reg *
self.params[str('W' + str(i+1))]

    #
===== #
    # END YOUR CODE HERE
    #
===== #

    return loss, grads

```