

```

# cnn .py

import numpy as np

from nndl.layers import *
from nndl.conv_layers import *
from utils.fast_layers import *
from nndl.layer_utils import *
from nndl.conv_layer_utils import *

import pdb

class ThreeLayerConvNet(object):
    """
    A three-layer convolutional network with the following architecture:

    conv - relu - 2x2 max pool - affine - relu - affine - softmax

    The network operates on minibatches of data that have shape (N, C,
    H, W)
    consisting of N images, each with height H and width W and with C
    input
    channels.
    """

    def __init__(self, input_dim=(3, 32, 32), num_filters=32,
filter_size=7,
                hidden_dim=100, num_classes=10, weight_scale=1e-3,
reg=0.0,
                dtype=np.float32, use_batchnorm=False):
        """
        Initialize a new network.

        Inputs:
        - input_dim: Tuple (C, H, W) giving size of input data
        - num_filters: Number of filters to use in the convolutional layer
        - filter_size: Size of filters to use in the convolutional layer
        - hidden_dim: Number of units to use in the fully-connected hidden
layer
        - num_classes: Number of scores to produce from the final affine
layer.
        - weight_scale: Scalar giving standard deviation for random
initialization
of weights.
        - reg: Scalar giving L2 regularization strength
        - dtype: numpy datatype to use for computation.
        """
        self.use_batchnorm = use_batchnorm
        self.params = {}
        self.reg = reg

```

```

self.dtype = dtype

# =====
#
# YOUR CODE HERE:
#   Initialize the weights and biases of a three layer CNN. To
initialize:
#       - the biases should be initialized to zeros.
#       - the weights should be initialized to a matrix with entries
#           drawn from a Gaussian distribution with zero mean and
#           standard deviation given by weight_scale.
# =====
#

C, H, W = input_dim

# cnn params
pad = (filter_size - 1) / 2
conv_stride = 1
pool_size = 2
pool_stride = 2
# output sizes after convolution and pooling
H_out_conv, W_out_conv = int(1 + (H - filter_size + 2*pad) /
conv_stride), int(1 + (W - filter_size + 2*pad) / conv_stride)
H_out_pool, W_out_pool = int(1 + (H_out_conv - pool_size) /
pool_stride), int(1 + (W_out_conv - pool_size) / pool_stride)

# W1 = conv weights
self.params['W1'] = weight_scale * np.random.randn(num_filters, C,
filter_size, filter_size)
self.params['b1'] = np.zeros(num_filters)

max_pool_output_size = int(num_filters * H_out_pool * W_out_pool)
self.params['W2'] = weight_scale *
np.random.randn(max_pool_output_size, hidden_dim)
self.params['b2'] = np.zeros(hidden_dim)
self.params['W3'] = weight_scale * np.random.randn(hidden_dim,
num_classes)
self.params['b3'] = np.zeros(num_classes)

# =====
#
# END YOUR CODE HERE
# =====
#

for k, v in self.params.items():
    self.params[k] = v.astype(dtype)

```

```

def loss(self, X, y=None):
    """
    Evaluate loss and gradient for the three-layer convolutional
    network.

    Input / output: Same API as TwoLayerNet in fc_net.py.
    """
    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']
    W3, b3 = self.params['W3'], self.params['b3']

    # pass conv_param to the forward pass for the convolutional layer
    filter_size = W1.shape[2]
    conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}

    # pass pool_param to the forward pass for the max-pooling layer
    pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

    scores = None

    # =====
    #
    # YOUR CODE HERE:
    #   Implement the forward pass of the three layer CNN. Store the
    output
    #   scores as the variable "scores".
    # =====
    #

    # conv - relu - 2x2 max pool - affine - relu - affine - softmax

    # get scores for first layer (conv + relu + pool)
    h1, cache1 = conv_relu_pool_forward(x=X, w=W1, b=b1,
conv_param=conv_param, pool_param=pool_param)
    # get scores for second layer (fc)
    h2, cache2 = affine_relu_forward(x=h1, w=W2, b=b2) # get scores
    for output layer (fc)
    scores, cache3 = affine_forward(x=h2, w=W3, b=b3)

    # =====
    #
    # END YOUR CODE HERE
    # =====
    #

    if y is None:
        return scores

    loss, grads = 0, {}

```

```

# =====
#
# YOUR CODE HERE:
#   Implement the backward pass of the three layer CNN. Store the
grads
#   in the grads dictionary, exactly as before (i.e., the gradient
of
#   self.params[k] will be grads[k]). Store the loss as "loss",
and
#   don't forget to add regularization on ALL weight matrices.
# =====
#

    loss, dl = softmax_loss(x=scores, y=y)
    loss += 0.5*self.reg*np.sum(W1**2) + 0.5*self.reg*np.sum(W2**2) +
0.5*self.reg*np.sum(W3**2)
    dout, dW3, db3 = affine_backward(dl, cache3) # now backprop,
starting from the last affine layer
    dW3 += self.reg * W3
    dout, dW2, db2 = affine_relu_backward(dout, cache2)
    dW2 += self.reg * W2
    dx, dW1, db1 = conv_relu_pool_backward(dout, cache1)
    dW1 += self.reg * W1
    # now store all the gradients in the gradient dictionary
    grads["W1"] = dW1
    grads["W2"] = dW2
    grads["W3"] = dW3
    grads["b1"] = db1
    grads["b2"] = db2
    grads["b3"] = db3

# =====
#
# END YOUR CODE HERE
# =====
#

```

```

    return loss, grads

```

```

class BestCNN(object):
    def __init__(self, input_dim=(3, 32, 32), num_filters=32,
filter_size=7, hidden_dim=100, num_classes=10, weight_scale=1e-3,
reg=0.0, dtype=np.float32, use_batchnorm=False):
    """
    Initialize a new network.
    Inputs:
    - input_dim: Tuple (C, H, W) giving size of input data
    - num_filters: Number of filters to use in the convolutional layer
    - filter_size: Size of filters to use in the convolutional layer

```

```

        - hidden_dim: Number of units to use in the fully-connected hidden
layer
        - num_classes: Number of scores to produce from the final affine
layer.
        - weight_scale: Scalar giving standard deviation for random
initialization
of weights.
        - reg: Scalar giving L2 regularization strength
        - dtype: numpy datatype to use for computation. ""
        self.use_batchnorm = use_batchnorm
        self.params = {}
        self.reg = reg
        self.dtype = dtype
        # =====
# # YOUR CODE HERE:
# # # # # #

        # plan
        # {conv relu conv relu pool} x 2 -> affine -> relu -> affine ->
output
        # bn plan
        # {conv bn relu conv bn relu pool} x 2 -> affine -> bn -> relu ->
affine -> output, thus need 5 bns

        C, H, W = input_dim
        # hyperparams to use
        pad = (filter_size - 1) / 2
        conv_stride = 1
        pool_size = 2
        pool_stride = 2

        # init
        self.params["W1"] = np.random.normal(loc=0, scale=weight_scale,
size=(num_filters, C, filter_size, filter_size))
        self.params["b1"] = np.zeros(num_filters)
        self.params["W2"] = np.random.normal(loc=0, scale=weight_scale,
size=(num_filters, num_filters, filter_size, filter_size))
        self.params["b2"] = np.zeros(num_filters)
        self.params["W3"] = np.random.normal(loc=0, scale=weight_scale,
size=(num_filters, num_filters, filter_size, filter_size))
        self.params["b3"] = np.zeros(num_filters)
        self.params["W4"] = np.random.normal(loc=0, scale=weight_scale,
size=(num_filters, num_filters, filter_size, filter_size))
        self.params["b4"] = np.zeros(num_filters)
        self.params["W5"] = np.random.normal(loc=0, scale=weight_scale,
size=(num_filters*8*8, hidden_dim)) ## 8 because this is the
H_out_pool
        self.params["b5"] = np.zeros(hidden_dim)

        # output layer is different

```

```

        self.params["W6"] = np.random.normal(loc=0, scale=weight_scale,
size=(hidden_dim, num_classes))
        self.params["b6"] = np.zeros(num_classes)

    if self.use_batchnorm:
        for i in range(1,5):
            self.params['gamma'+str(i)] = np.ones(num_filters)
            self.params['beta'+str(i)] = np.zeros(num_filters)

        self.params['gamma5'] = np.ones(hidden_dim)
        self.params['beta5'] = np.zeros(hidden_dim)

        self.bn_params = []
        if self.use_batchnorm:
            self.bn_params = [{'mode': 'train'} for i in np.arange(5)]

    # =====
    # # END YOUR CODE HERE
    # =====
    #
    for k, v in self.params.items(): self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    # print("input shape", X.shape)
    mode = 'test' if y is None else 'train'

    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']
    W3, b3 = self.params['W3'], self.params['b3']
    W4, b4 = self.params['W4'], self.params['b4']
    W5, b5 = self.params['W5'], self.params['b5']
    W6, b6 = self.params['W6'], self.params['b6']

    if self.use_batchnorm:
        for bn_param in self.bn_params:
            bn_param['mode'] = mode

    # take care of conv
    filter_size = W1.shape[2]
    conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}
    # take care of pool
    pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}
    scores = None

    if not self.use_batchnorm:
        # forward pass w/o bn
        h1, cache1 = conv_relu_forward(x=X, w=W1, b=b1,

```

```

conv_param=conv_param) # conv relu

    h2, cache2 = conv_relu_forward(x=h1, w=W2, b=b2,
conv_param=conv_param) # conv relu
    h3, cache3 = max_pool_forward_fast(x=h2, pool_param=pool_param)
# pool
    h4, cache4 = conv_relu_forward(x=h3, w=W3, b=b3,
conv_param=conv_param) # conv relu
    h5, cache5 = conv_relu_forward(x=h4, w=W4, b=b4,
conv_param=conv_param) # conv relu
    h6, cache6 = max_pool_forward_fast(x=h5, pool_param=pool_param)
# pool
    # FC
    h7, cache7 = affine_relu_forward(x=h6, w=W5, b=b5) # affine
    scores, cache8 = affine_forward(x=h7, w=W6, b=b6) # affine -
output

```

```

bn_cache=[]
# forward pass w bn
if self.use_batchnorm:
    h1, cache1 = conv_bn_relu_forward(x=X, w=W1, b=b1,
conv_param=conv_param, gamma=self.params['gamma1'],
beta=self.params['beta1'], bn_param=self.bn_params[0]) # conv relu
    h2, cache2 = conv_bn_relu_forward(x=h1, w=W2, b=b2,
conv_param=conv_param, gamma=self.params['gamma2'],
beta=self.params['beta2'], bn_param=self.bn_params[1]) # conv relu
    h3, cache3 = max_pool_forward_fast(x=h2, pool_param=pool_param)
# pool
    h4, cache4 = conv_bn_relu_forward(x=h3, w=W3, b=b3,
conv_param=conv_param, gamma=self.params['gamma3'],
beta=self.params['beta3'], bn_param=self.bn_params[2]) # conv relu
    h5, cache5 = conv_bn_relu_forward(x=h4, w=W4, b=b4,
conv_param=conv_param, gamma=self.params['gamma4'],
beta=self.params['beta4'], bn_param=self.bn_params[3]) # conv relu
    h6, cache6 = max_pool_forward_fast(x=h5, pool_param=pool_param)
# pool
    # FC
    h7, cache7 = affine_batchnorm_relu_forward(x=h6, w=W5, b=b5,
gamma=self.params['gamma5'], beta=self.params['beta5'],
bn_params=self.bn_params[4]) # affine
    scores, cache8 = affine_forward(x=h7, w=W6, b=b6) # affine -
output

```

```

if y is None:
    return scores
loss, grads = 0, {}

```

```

    loss, dl = softmax_loss(x=scores, y=y) # then regularize the loss
    loss += 0.5*self.reg*np.sum(W1**2) + 0.5*self.reg*np.sum(W2**2) +
0.5*self.reg*np.sum(W3**2) + 0.5*self.reg*np.sum(W4**2) +
0.5*self.reg*np.sum(W5**2) + 0.5*self.reg*np.sum(W6**2)

    if not self.use_batchnorm:
        dout, dW6, db6 = affine_backward(dl, cache8) # affine
        dW6 += self.reg * W6
        dout, dW5, db5 = affine_relu_backward(dout, cache7) # affine
relu
        dW5 += self.reg * W5

        dout = max_pool_backward_fast(dout, cache6) # pool
        dout, dW4, db4 = conv_relu_backward(dout, cache5) # conv relu
        dW4 += self.reg * W4
        dout, dW3, db3 = conv_relu_backward(dout, cache4) # conv relu
        dW3 += self.reg * W3
        dout = max_pool_backward_fast(dout, cache3) # pool
        dout, dW2, db2 = conv_relu_backward(dout, cache2) # conv relu
        dW2 += self.reg * W2
        dout, dW1, db1 = conv_relu_backward(dout, cache1) # conv relu
        dW1 += self.reg * W1

    else:
        dout, dW6, db6 = affine_backward(dl, cache8) # affine
        dW6 += self.reg * W6
        dout, dW5, db5, dgamma5, dbeta5 =
affine_batchnorm_relu_backward(dout, cache7) # affine relu
        dW5 += self.reg * W5
        grads['gamma5'] = dgamma5
        grads['beta5'] = dbeta5

        dout = max_pool_backward_fast(dout, cache6) # pool
        dout, dW4, db4, dgamma4, dbeta4 = conv_bn_relu_backward(dout,
cache5) # conv relu
        dW4 += self.reg * W4
        grads['gamma4'] = dgamma4
        grads['beta4'] = dbeta4

        dout, dW3, db3, dgamma3, dbeta3 = conv_bn_relu_backward(dout,
cache4) # conv relu
        grads['gamma3'] = dgamma3
        grads['beta3'] = dbeta3
        dW3 += self.reg * W3

        dout = max_pool_backward_fast(dout, cache3) # pool
        dout, dW2, db2, dgamma2, dbeta2 = conv_bn_relu_backward(dout,
cache2) # conv relu

```



```

        grads['gamma2'] = dgamma2
        grads['beta2'] = dbeta2
        dW2 += self.reg * W2
        dout, dW1, db1, dgamma1, dbeta1 = conv_bn_relu_backward(dout,
cache1) # conv relu
        grads['gamma1'] = dgamma1
        grads['beta1'] = dbeta1
        dW1 += self.reg * W1

```

```

# storage of w's and b's
grads["W1"] = dW1
grads["W2"] = dW2
grads["W3"] = dW3
grads["W4"] = dW4
grads["W5"] = dW5
grads["W6"] = dW6
grads["b1"] = db1
grads["b2"] = db2
grads["b3"] = db3
grads["b4"] = db4
grads["b5"] = db5
grads["b6"] = db6

```

```

return loss, grads

```