

```

# conv layers updated

import numpy as np
from nndl.layers import *
import pdb

def conv_forward_naive(x, w, b, conv_param):
    """
    A naive implementation of the forward pass for a convolutional
    layer.

    The input consists of N data points, each with C channels, height H
    and width W. We convolve each input with F different filters, where each
    filter spans
    all C channels and has height HH and width HH.

    Input:
    - x: Input data of shape (N, C, H, W)
    - w: Filter weights of shape (F, C, HH, WW)
    - b: Biases, of shape (F,)
    - conv_param: A dictionary with the following keys:
        - 'stride': The number of pixels between adjacent receptive fields
        in the horizontal and vertical directions.
        - 'pad': The number of pixels that will be used to zero-pad the
        input.

    Returns a tuple of:
    - out: Output data, of shape (N, F, H', W') where H' and W' are
    given by
         $H' = 1 + (H + 2 * \text{pad} - \text{HH}) / \text{stride}$ 
         $W' = 1 + (W + 2 * \text{pad} - \text{WW}) / \text{stride}$ 
    - cache: (x, w, b, conv_param)
    """
    out = None
    pad = conv_param['pad']
    stride = conv_param['stride']

    # ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of a convolutional neural network.
    # Store the output as 'out'.
    # Hint: to pad the array, you can use the function np.pad.
    # ===== #
    N, C, H, W = x.shape
    F, C, HH, WW = w.shape

    # check if valid conv

```

```

    assert (H + 2 * pad - HH) % stride == 0
    assert (W + 2 * pad - WW) % stride == 0

    # only pad the third and fourth axes
    x_pad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)),
mode='constant')

    # H_prime, W_prime are the output height and width
    H_prime = int(1 + (H + 2 * pad - HH) / stride)
    W_prime = int(1 + (W + 2 * pad - WW) / stride)

    out = np.zeros((N, F, H_prime, W_prime))

    for n in range(N): # number of batches
        for i in range(H_prime): # output height
            for j in range(W_prime): # output width
                seg = x_pad[n,:,i*stride:i*stride + HH, j*stride:j*stride +
WW]
                out[n,:,i,j] = np.sum(seg * w,axis=(1,2,3)) + b

# ===== #
# END YOUR CODE HERE
# ===== #

    cache = (x, w, b, conv_param)
    return out, cache

def conv_backward_naive(dout, cache):
    """
    A naive implementation of the backward pass for a convolutional
    layer.

    Inputs:
    - dout: Upstream derivatives.
    - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive

    Returns a tuple of:
    - dx: Gradient with respect to x
    - dw: Gradient with respect to w
    - db: Gradient with respect to b
    """
    dx, dw, db = None, None, None

    N, F, out_height, out_width = dout.shape
    x, w, b, conv_param = cache

    stride, pad = [conv_param['stride'], conv_param['pad']]

```

```

xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)),
mode='constant')
num_filt, _, f_height, f_width = w.shape

# ===== #
# YOUR CODE HERE:
#   Implement the backward pass of a convolutional neural network.
#   Calculate the gradients: dx, dw, and db.
# ===== #
# inits
dx = np.zeros_like(x)
dw = np.zeros_like(w)
db = np.zeros_like(b)

# dims
N, C, H, W = x.shape
F, _, HH, WW = w.shape
_, _, H_out, W_out = dout.shape

# db
db = np.sum(dout, axis=(0, 2, 3))

# dw
for f in range(F): # looping through filters
    for c in range(C): # looping through channels
        for i in range(HH): # looping through weight height
            for j in range(WW): # looping through weight width
                dw[f, c, i, j] = np.sum(xpad[:, c, i: i + H_out *
stride : stride, j : j + W_out* stride : stride] * dout[:, f, :, :])
                # dw at determined segment
                # np sum of (xpad matrix * dout matrix)

# dx
dx = np.zeros(x.shape)
# loop through the number of examples
for n in range(N): # hi, wi: loop through x
    for hx in range(H):
        for wx in range(W):
            y_indexes = [] # will contain valid indices of y
            w_indexes = [] # will contain valid indices of weight
(w)
            for i in range(H_out): # H_ is from dout
                for j in range(W_out): # W_ is from dout
                    # i, j: loop through output
                    # verify: is the range within weights limits?
                    h_range = (hx + pad - i * stride) # height range
                    w_range = (wx + pad - j * stride) # weight range
                    if (h_range >= 0) and (h_range < HH) and
(w_range >= 0) and (w_range < WW):

```

```

        w_indexes.append((h_range, w_range))
        y_indexes.append((i, j))

    for f in range(F): # filters loop
        # windex_f and yindex_f from python zip of
w_indexes, y_indexes (as determined above)
        # increment by np.sum ( w_matrix * dout_matrix) for
valid indices of y and weights
        dx[n, : , hx, wx] += np.sum([w[f, :, windex_f[0],
windex_f[1]] * dout[n, f, yindex_f[0], yindex_f[1]] for windex_f,
yindex_f in zip(w_indexes, y_indexes)], 0)

# ===== #
# END YOUR CODE HERE
# ===== #

return dx, dw, db

def max_pool_forward_naive(x, pool_param):
    """
    A naive implementation of the forward pass for a max pooling layer.

    Inputs:
    - x: Input data, of shape (N, C, H, W)
    - pool_param: dictionary with the following keys:
        - 'pool_height': The height of each pooling region
        - 'pool_width': The width of each pooling region
        - 'stride': The distance between adjacent pooling regions

    Returns a tuple of:
    - out: Output data
    - cache: (x, pool_param)
    """
    out = None

    # ===== #
    # YOUR CODE HERE:
    #   Implement the max pooling forward pass.
    # ===== #

    N, C, H, W = x.shape
    HH = pool_param['pool_height']
    WW = pool_param['pool_width']
    stride = pool_param['stride']

    # H_prime, W_prime are the output height and width
    H_prime = int(1 + (H - HH) / stride)

```

```

W_prime = int(1 + (W - WW) / stride)

out = np.zeros((N, C, H_prime, W_prime))

for n in range(N): # number of batches
    for i in range(H_prime): # output height
        for j in range(W_prime): # output width
            seg = x[n,:,i*stride:i*stride + HH, j*stride:j*stride + WW]
            out[n,:,i,j] = np.amax(seg,axis=(1,2))

# ===== #
# END YOUR CODE HERE
# ===== #
cache = (x, pool_param)
return out, cache

def max_pool_backward_naive(dout, cache):
    """
    A naive implementation of the backward pass for a max pooling layer.

    Inputs:
    - dout: Upstream derivatives
    - cache: A tuple of (x, pool_param) as in the forward pass.

    Returns:
    - dx: Gradient with respect to x
    """
    dx = None
    x, pool_param = cache
    pool_height, pool_width, stride = pool_param['pool_height'],
    pool_param['pool_width'], pool_param['stride']

    # ===== #
    # YOUR CODE HERE:
    # Implement the max pooling backward pass.
    # ===== #

    N, C, H, W = x.shape
    N, C, dout_height, dout_width = dout.shape

    dx = np.zeros_like(x)

    for n in range(N): # loop over the
        number of training samples
        for c in range(C): # loop over the number
            of channels
                for i in range(dout_height): # loop over vertical
                    axis of the dout
                        for j in range(dout_width): # loop over horizontal
                            axis of the dout

```

```

        i_, j_ = np.where(np.max(x[n, c, i * stride : i *
stride + pool_height, j * stride : j * stride + pool_width]) == x[n,
c, i * stride : i * stride + pool_height, j * stride : j * stride +
pool_width])
        i_, j_ = i_[0], j_[0]
        dx[n, c, i * stride : i * stride + pool_height, j *
stride : j * stride + pool_width][i_, j_] = dout[n, c, i, j]

# ===== #
# END YOUR CODE HERE
# ===== #

return dx

def spatial_batchnorm_forward(x, gamma, beta, bn_param):
    """
    Computes the forward pass for spatial batch normalization.

    Inputs:
    - x: Input data of shape (N, C, H, W)
    - gamma: Scale parameter, of shape (C,)
    - beta: Shift parameter, of shape (C,)
    - bn_param: Dictionary with the following keys:
        - mode: 'train' or 'test'; required
        - eps: Constant for numeric stability
        - momentum: Constant for running mean / variance. momentum=0 means
that
old information is discarded completely at every time step,
while
momentum=1 means that new information is never incorporated. The
default of momentum=0.9 should work well in most situations.
        - running_mean: Array of shape (D,) giving running mean of
features
        - running_var Array of shape (D,) giving running variance of
features

    Returns a tuple of:
    - out: Output data, of shape (N, C, H, W)
    - cache: Values needed for the backward pass
    """
    out, cache = None, None

# ===== #
# YOUR CODE HERE:
# Implement the spatial batchnorm forward pass.
#
# You may find it useful to use the batchnorm forward pass you
# implemented in HW #4.
# ===== #

```

```

N, C, W, H = x.shape
xr = x.reshape(N*H*W, C)
out, cache = batchnorm_forward(xr, gamma, beta, bn_param)
out = out.reshape(N, C, W, H)

# ===== #
# END YOUR CODE HERE
# ===== #

return out, cache

def spatial_batchnorm_backward(dout, cache):
    """
    Computes the backward pass for spatial batch normalization.

    Inputs:
    - dout: Upstream derivatives, of shape (N, C, H, W)
    - cache: Values from the forward pass

    Returns a tuple of:
    - dx: Gradient with respect to inputs, of shape (N, C, H, W)
    - dgamma: Gradient with respect to scale parameter, of shape (C,)
    - dbeta: Gradient with respect to shift parameter, of shape (C,)
    """
    dx, dgamma, dbeta = None, None, None

    # ===== #
    # YOUR CODE HERE:
    #   Implement the spatial batchnorm backward pass.
    #
    #   You may find it useful to use the batchnorm forward pass you
    #   implemented in HW #4.
    # ===== #

    N, C, W, H = dout.shape

    doutr = dout.reshape(N*H*W, C)

    dx, dgamma, dbeta = batchnorm_backward(doutr, cache)

    dx = dx.reshape(N, C, W, H)
    dgamma = dgamma.reshape(C,)
    dbeta = dbeta.reshape(C,)

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return dx, dgamma, dbeta

```